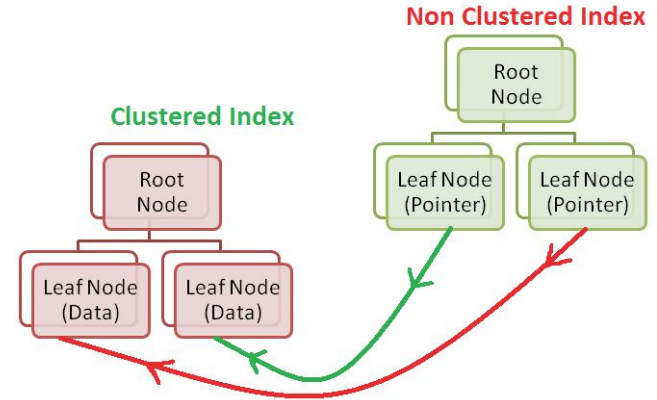


# Clustered Index

- Storing the indexed row directly within an index
  - Telling the order of the data in the index
- Why?
  - Reads performance can be bad, when
    - the location of the heap file keeps changed, needing additional pointers
- Only one clustered index per table in SQL
  - In some system, primary key is always the clustered index
- Adds overhead on write requiring additional storage and consistency



# Concatenated Index: Multi-column Index

- Combines several fields into one key
  - Why?
    - A single key is not enough when querying multiple columns simultaneously
  - Examples
    - Geospatial data - latitude, longitude
    - Weather observation - date, temperature
    - Selecting colours for product search - red/green/blue
  - B-tree or LSM-tree can't handle these as it is
    - Needs translation process to R-tree
-

### Full-text search and fuzzy indexes

- Search for similar keys
  - misspelled words, synonyms
- Lucene
  - For the term dictionary, SSTable-like structure with a small in-memory index is used to find the offset
- Techniques for Machine learning

### In-memory database

- RAM becomes cheaper
- Big performance improvements claimed
  - Avoiding encoding in-memory data structures to be written in a disk
- Different data models from disk-based index
  - Priority queues & sets of Redis

---

# Data Warehousing

Laine Kim

*From <Designing Data-Intensive Applications>  
by Martin Kleppmann*

---

# OLTP

## Online Transaction Processing

- End user, customers
- Read: limited records per query
- Write: random, low-latency from user inputs
- Highly available

# OLAP

## Online Analytic Processing

- Used for business analysis
  - Read: Aggregate numerous historical records
  - Write: Bulk import
  - Larger dataset size than OLTP
-

# Data Warehouse

- **Separate Database** from the OLTP system
  - Read-only copy of the OLTP
  - Optimized for analytic specific access patterns
  - Writing is more difficult but LSM-trees is a solution
    - Not real-time write
    - First write in-memory store, being sorted and prepared for writing to disk until enough writes accumulated
  - Most Commonly relational as SQL
-

# Star schema for Analytics




Figure 3-9 Example of a star schema for use in a data warehouse,  
<Designing Data-Intensive Applications>

# Row-Oriented Storage



- Each row is stored next to each other
- Slow to fetch a million rows for analysis
- Applicable to OLTP

# Column-Oriented Storage



- Store values from each column together
  - Rows in the same order
- Efficient use of CPU cycle
  - Compressed column data allows more rows to fit in the CPU's L1 cache



# Column Compression - bitmap encoding

Column values:

product\_sk: 69 69 69 69 74 31 31 31 29 30 30 31 31 31 68 69 69

Bitmap for each value

product\_sk = 31: 0 0 0 0 0 1 1 1 1 0 0 0 1 1 1 0 0 0

product\_sk = 68: 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0

product\_sk = 69: 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1

Run-length encoding

product\_sk = 31: 5, 4, 3, 3 (5 zeros, 4 ones, 3 zeros, 3 ones, rest zeros)

product\_sk = 68: 15, 1 (15 zeros, 1 one, rest zeros)

product\_sk = 69: 0, 4, 12, 2 (0 zeros, 4 ones, 12 zeros, 2 ones)

- Example query

- WHERE product\_sk IN (30, 68, 69):
- Load the three bitmaps with OR operation, efficiently!

A row

Figure 3-11. Compressed, bitmap-indexed storage of a single column,  
<Designing Data-Intensive Applications>

# Sort Order

- Sort an entire row at a time
  - Otherwise, can't reconstruct the row
- Set sort key to be the most frequently used column
  - Ex) often targeting date ranges, *data\_key* is the first sort key
- Sorted order helps the compression of the columns

# Aggregation with Materialized Views

- Cache often used queries with aggregate functions
    - COUNT, SUM, AVG, MIN, or MAX in SQL
  - Needs to be updated as data changes
  - Precomputed, thus very fast
  - Less complexity than raw data - limited usages
-

---

# The End

---